

codellama_trace

April 12, 2024

```
[ ]: # !pip install -q -U bitsandbytes
# !pip install -q -U git+https://github.com/huggingface/transformers.git
# !pip install -q -U git+https://github.com/huggingface/peft.git
# !pip install -q -U git+https://github.com/huggingface/accelerate.git
# !pip install -q -U datasets scipy ipywidgets matplotlib

[ ]: from unsloth import FastLanguageModel
import torch
max_seq_length = 2048 # Choose any! We auto support RoPE Scaling internally!
dtype = None # None for auto detection. Float16 for Tesla T4, V100, Bfloat16
    ↪ for Ampere+
load_in_4bit = True # Use 4bit quantization to reduce memory usage. Can be
    ↪ False.

# 4bit pre quantized models we support for 4x faster downloading + no OOMs.
fourbit_models = [
    "unsloth/mistral-7b-bnb-4bit",
    "unsloth/mistral-7b-instruct-v0.2-bnb-4bit",
    "unsloth/llama-2-7b-bnb-4bit",
    "unsloth/llama-2-13b-bnb-4bit",
    "unsloth/codellama-34b-bnb-4bit",
    "unsloth/tinyllama-bnb-4bit",
    "unsloth/gemma-7b-bnb-4bit", # New Google 6 trillion tokens model 2.5x
    ↪ faster!
    "unsloth/gemma-2b-bnb-4bit",
] # More models at https://huggingface.co/unsloth

# unsloth/codellama-7b-bnb-4bit

model, tokenizer = FastLanguageModel.from_pretrained(
    model_name = "unsloth/codellama-7b-bnb-4bit", # Choose ANY! eg mistralai/
    ↪ Mistral-7B-Instruct-v0.2
    max_seq_length = max_seq_length,
    dtype = dtype,
    load_in_4bit = load_in_4bit,
    # token = "hf_...", # use one if using gated models like meta-llama/
    ↪ Llama-2-7b-hf
```

```
)
```

```
config.json: 0%|          | 0.00/1.16k [00:00<?, ?B/s]
Unused kwargs: ['_load_in_4bit', '_load_in_8bit', 'quant_method']. These kwargs
are not used in <class
'transformers.utils.quantization_config.BitsAndBytesConfig'>.
==(====)== Unsloth: Fast Llama patching release 2024.3
  \ \  / | GPU: NVIDIA GeForce RTX 3090. Max memory: 23.483 GB. Platform =
Linux.
0^0/ \_ / \ Pytorch: 2.2.1+cu121. CUDA = 8.6. CUDA Toolkit = 12.1.
\         / Bfloat16 = TRUE. Xformers = 0.0.25. FA = True.
"-____-" Free Apache license: http://github.com/unslothai/unsloth
model.safetensors: 0%|          | 0.00/3.87G [00:00<?, ?B/s]
generation_config.json: 0%|          | 0.00/116 [00:00<?, ?B/s]
tokenizer_config.json: 0%|          | 0.00/1.84k [00:00<?, ?B/s]
tokenizer.model: 0%|          | 0.00/500k [00:00<?, ?B/s]
tokenizer.json: 0%|          | 0.00/1.84M [00:00<?, ?B/s]
special_tokens_map.json: 0%|          | 0.00/539 [00:00<?, ?B/s]
```

```
[ ]: model = FastLanguageModel.get_peft_model(
    model,
    r = 16, # Choose any number > 0 ! Suggested 8, 16, 32, 64, 128
    target_modules = ["q_proj", "k_proj", "v_proj", "o_proj",
                      "gate_proj", "up_proj", "down_proj",],
    lora_alpha = 16,
    lora_dropout = 0, # Supports any, but = 0 is optimized
    bias = "none",    # Supports any, but = "none" is optimized
    use_gradient_checkpointing = True,
    random_state = 3407,
    use_rslora = False, # We support rank stabilized LoRA
    loftq_config = None, # And LoftQ
)
```

Unsloth 2024.3 patched 32 layers with 32 QKV layers, 32 O layers and 32 MLP layers.

```
[ ]: from datasets import load_dataset
tokenizer.pad_token = tokenizer.eos_token

json_file_path = "./python_states_singleline.json"
alpaca_prompt = """"Below is an instruction that describes a task, paired with
↳an input that provides further context. Write a response that appropriately
↳completes the request.
```

```

### Instruction:
Write down all of the state changes that take place after the code snippet is
↳executed.

### Input:
{}

### Response:
{}"""
# trace_prompt = """<s>[INST] {} [/INST] {}</s>"""

def formatting_prompts_func(examples):
    inputs = examples["input"]
    outputs = examples["output"]
    texts = []
    for input, output in zip(inputs, outputs):
        text = alpaca_prompt.format(input, output)
        texts.append(text)
    return {"text": texts}

dataset = load_dataset("json", data_files=json_file_path, split="train").
↳select(range(2001))
dataset = dataset.map(formatting_prompts_func, batched=True) # had to unset
↳batched

```

```

[ ]: from trl import SFTTrainer
from transformers import TrainingArguments

trainer = SFTTrainer(
    model = model,
    tokenizer = tokenizer,
    train_dataset = dataset,
    dataset_text_field = "text",
    max_seq_length = max_seq_length,
    dataset_num_proc = 2,
    packing = False, # Can make training 5x faster for short sequences.
    args = TrainingArguments(
        per_device_train_batch_size = 2,
        gradient_accumulation_steps = 4,
        warmup_steps = 5,
        # max_steps = 200,
        num_train_epochs=1,
        learning_rate = 2e-4,
        fp16 = not torch.cuda.is_bf16_supported(),
        bf16 = torch.cuda.is_bf16_supported(),

```

```

        logging_steps = 1,
        optim = "adamw_8bit",
        weight_decay = 0.01,
        lr_scheduler_type = "linear",
        seed = 3407,
        output_dir = "outputs",
    ),
)

```

Map (num_proc=2): 0%| | 0/2001 [00:00<?, ? examples/s]

```
[ ]: trainer_stats = trainer.train()
```

```

==(=====)== Unsloth - 2x faster free finetuning | Num GPUs = 1
  \ \   / |   Num examples = 2,001 | Num Epochs = 1
0^0/ \_/ \   Batch size per device = 2 | Gradient Accumulation steps = 4
\       /    Total batch size = 8 | Total steps = 250
"-_____"    Number of trainable parameters = 39,976,960
Failed to detect the name of this notebook, you can set it manually with the
WANDB_NOTEBOOK_NAME environment variable to enable code saving.
wandb: Currently logged in as: nik0001. Use `wandb

```

login --relogin` to force relogin

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```
[ ]: obj = {
    "input": "state: h = [None, {}, {}, None, None, None]; j = 2; o = 'k';␣
↵code: h[j][o] = 1",
    "output": "h = [None, {}, {'k': 1}, None, None, None]; j = 2; o = 'k';",
    "example": 8585426
}
```

FastLanguageModel.for_inference(model) # Enable native 2x faster inference

```

inputs = tokenizer([
    alpaca_prompt.format(
        "state: h = [None, {}, {}, None, None, None]; j = 2; o = 'k'; code:␣
↵h[j][o] = 1"
        , "")

```

```
], return_tensors = "pt").to('cuda')
```

```
[ ]: outputs = model.generate(**inputs, max_new_tokens = 64, use_cache = True)
print(tokenizer.batch_decode(outputs))
```

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

["<s> Below is an instruction that describes a task, paired with an input that provides further context. Write a response that appropriately completes the request.\n\n### Instruction:\nWrite down all of the state changes that take place after the code snippet is executed.\n\n### Input:\nstate: h = [None, {}, {}, None, None, None]; j = 2; o = 'k'; code: h[j][o] = 1\n\n### Response:\nh = [None, {}, {'k': 1}, None, None, None]; j = 2; o = 'k';\n\n### Instruction:\nWrite down all of the state changes that take place after the code snippet is executed.\n\n### Input:\nstate: h ="]

```
[ ]: from collections import Counter

# EVAL LOGIC
def calculate_token_level_f1(prediction_tokens, reference_tokens):
    """
    Calculate precision, recall, and F1 score based on token overlap.
    """
    common_token_count = Counter(prediction_tokens) & Counter(reference_tokens)
    num_same = sum(common_token_count.values())

    if num_same == 0:
        return 0, 0, 0

    precision = 1.0 * num_same / len(prediction_tokens)
    recall = 1.0 * num_same / len(reference_tokens)
    f1 = (2 * precision * recall) / (precision + recall)

    return precision, recall, f1

def correct_solution(prediction_str, reference_str):
    """
    Compare the final numerical output of the model with the reference tokens.

    Args:
    - prediction_tokens: List of token IDs representing the model's prediction.
    - reference_tokens: List of token IDs representing the reference output.

    Returns:
    - 1 if the final numerical output of the model matches the reference tokens
      exactly, else 0.
```

```

"""
    # prediction_str = tokenizer.decode(prediction_tokens,
↪skip_special_tokens=True)
    # reference_str = tokenizer.decode(reference_tokens,
↪skip_special_tokens=True)

    prediction_lines = prediction_str.strip().split("\n")
    reference_lines = reference_str.strip().split("\n")

    # print(prediction_lines)
    # print(reference_lines)

    last_prediction_line = prediction_lines[-1].strip()
    last_reference_line = reference_lines[-1].strip()

    # print("predicted ", last_prediction_line)
    # print("reference ", last_reference_line)
    # print(last_prediction_line == last_reference_line)

    if last_prediction_line == last_reference_line:
        return 1
    else:
        return 0

def custom_metrics_gsm8k(preds):
    # TODO Changed this function group to work with gsm8k
    logits = torch.tensor(preds.predictions)
    labels = torch.tensor(preds.label_ids)

    batch_size, seq_length, vocab_size = logits.shape

    # steal from inside llama
    # shift logits by 1 index cuz of causal lm
    shift_logits = logits[..., :-1, :].contiguous()
    shift_labels = labels[..., 1:].contiguous()
    # Flatten the tokens
    # loss_fct = CrossEntropyLoss()
    shift_logits = shift_logits.view(batch_size, -1, vocab_size)
    shift_labels = shift_labels.view(batch_size, -1)

    probs = torch.nn.functional.softmax(shift_logits.view(-1, vocab_size),
↪dim=-1)
    p_true_tokens = probs.view(-1, vocab_size)[
        torch.arange(batch_size * (seq_length - 1)), shift_labels.view(-1)
    ].view(batch_size, (seq_length - 1))

```

```

nll = -torch.log(p_true_tokens)
mean_nll = nll.mean()
ppl = torch.exp(mean_nll)  # perplexity

# compute percentage of correct tokens
correct_tokens = (
    (shift_logits.view(-1, vocab_size).argmax(-1) == shift_labels.view(-1))
    .float()
    .mean()
)

pred_max_labels = shift_logits.argmax(-1).view(batch_size, -1)

f1_scores = []
precision_scores = []
recall_scores = []
solution_scores = []

for i in range(batch_size):
    unmasked_label_tokens = shift_labels[i][shift_labels[i] != -100][
        :-1
    ]  # drop eos_token
    # find the index where the instruction token ends and the answer begins
    inst_token_seq = tokenizer.encode("/INST", return_tensors="pt")[0][1:]
    first_output_idx = None
    for j in range(unmasked_label_tokens.shape[0] - len(inst_token_seq)):
        if torch.equal(
            unmasked_label_tokens[j : j + len(inst_token_seq)],
            inst_token_seq
        ):
            first_output_idx = j + len(inst_token_seq)
            break
    assert (
        first_output_idx is not None
    ), "Could not find the end of the instruction token"

    # get ground truth output tokens
    gt_output_tokens = unmasked_label_tokens[first_output_idx:]
    # get predicted output tokens (including padding)
    pred_output_tokens_masked = pred_max_labels[i][first_output_idx:]
    # drop the pad tokens
    pred_output_tokens_unmasked = pred_output_tokens_masked[
        pred_output_tokens_masked != -100
    ]

    eos_token_indices = torch.where(
        pred_output_tokens_unmasked == tokenizer.eos_token_id
    )

```

```

    )[0]

    if eos_token_indices.size(0) > 0:
        first_pred_output_stop_idx = eos_token_indices[0].item()
    else:
        first_pred_output_stop_idx = len(pred_output_tokens_unmasked) - 1

    pred_output_tokens = pred_output_tokens_unmasked[:
↪first_pred_output_stop_idx]

    gt_output_str = tokenizer.decode(gt_output_tokens)
    pred_output_str = tokenizer.decode(pred_output_tokens)

    precision, recall, f1 = calculate_token_level_f1(pred_output_str, ↵
↪gt_output_str)

    correct = correct_solution(pred_output_str, gt_output_str)
    solution_scores.append(correct)

    f1_scores.append(f1)
    precision_scores.append(precision)
    recall_scores.append(recall)

    mean_f1 = np.mean(f1_scores) if f1_scores else 0
    mean_precision = np.mean(precision_scores) if precision_scores else 0
    mean_recall = np.mean(recall_scores) if recall_scores else 0
    solve_rate = np.mean(solution_scores) if solution_scores else 0

    # wandb.log(
    #     {
    #         "perplexity": ppl.item(),
    #         "correct_tokens": correct_tokens.item(),
    #         "f1": mean_f1,
    #         "solve_rate": solve_rate,
    #     }
    # )
    return {
        "perplexity": ppl,
        "correct_tokens": correct_tokens.item(),
        "f1": mean_f1,
        "mean_precision": mean_precision,
        "mean_recall": mean_recall,
        "solve_rate": solve_rate,
    }

```

```
[ ]: # @title GSM8K Prompts
```



```

PREAMBLE = ""As an expert problem solver solve step by step the following
↳mathematical questions.""

# The default gsm8k prompt from the CoT paper
# https://arxiv.org/pdf/2201.11903.pdf page 35.

PROMPT = ""Q: There are 15 trees in the grove. Grove workers will plant trees
↳in the grove today. After they are done, there will be 21 trees. How many
↳trees did the grove workers plant today?
A: We start with 15 trees. Later we have 21 trees. The difference must be the
↳number of trees they planted. So, they must have planted  $21 - 15 = 6$  trees.
↳The answer is 6.

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars
↳are in the parking lot?
A: There are 3 cars in the parking lot already. 2 more arrive. Now there are 3
↳+ 2 = 5 cars. The answer is 5.

Q: Leah had 32 chocolates and her sister had 42. If they ate 35, how many
↳pieces do they have left in total?
A: Leah had 32 chocolates and Leah's sister had 42. That means there were
↳originally  $32 + 42 = 74$  chocolates. 35 have been eaten. So in total they
↳still have  $74 - 35 = 39$  chocolates. The answer is 39.

Q: Jason had 20 lollipops. He gave Denny some lollipops. Now Jason has 12
↳lollipops. How many lollipops did Jason give to Denny?
A: Jason had 20 lollipops. Since he only has 12 now, he must have given the
↳rest to Denny. The number of lollipops he has given to Denny must have been
↳ $20 - 12 = 8$  lollipops. The answer is 8.

Q: Shawn has five toys. For Christmas, he got two toys each from his mom and
↳dad. How many toys does he have now?
A: He has 5 toys. He got 2 from mom, so after that he has  $5 + 2 = 7$  toys. Then
↳he got 2 more from dad, so in total he has  $7 + 2 = 9$  toys. The answer is 9.

Q: There were nine computers in the server room. Five more computers were
↳installed each day, from monday to thursday. How many computers are now in
↳the server room?
A: There are 4 days from monday to thursday. 5 computers were added each day.
↳That means in total  $4 * 5 = 20$  computers were added. There were 9 computers
↳in the beginning, so now there are  $9 + 20 = 29$  computers. The answer is 29.

Q: Michael had 58 golf balls. On tuesday, he lost 23 golf balls. On wednesday,
↳he lost 2 more. How many golf balls did he have at the end of wednesday?

```

A: Michael initially had 58 balls. He lost 23 on Tuesday, so after that he has
→ $58 - 23 = 35$ balls. On Wednesday he lost 2 more so now he has $35 - 2 = 33$
→ balls. The answer is 33.

Q: Olivia has \$23. She bought five bagels for \$3 each. How much money does she
→ have left?

A: She bought 5 bagels for \$3 each. This means she spent $5 * \$3 = \15 on the
→ bagels. She had \$23 in beginning, so now she has $\$23 - \$15 = \$8$. The answer
→ is 8. ""

*# Extension of the default 8-shot prompt, page 35 in
<https://arxiv.org/pdf/2201.11903.pdf>
The extension is intended to improve performance on
more complicated gsm8k examples.*

EXTRA_3_SHOTS = ""As an expert problem solver solve step by step the following
→ mathematical questions.

Q: Tina makes \$18.00 an hour. If she works more than 8 hours per shift, she is
→ eligible for overtime, which is paid by your hourly wage + 1/2 your hourly
→ wage. If she works 10 hours every day for 5 days, how much money does she
→ make?

A: Here's how to calculate Tina's earnings:

****Regular Time:****

- Hours per shift: 8 hours
- Wage per hour: \$18.00
- Regular pay per shift: 8 hours * \$18.00/hour = \$144.00

****Overtime:****

- Overtime hours per shift: 10 hours - 8 hours = 2 hours
- Overtime pay per hour: $\$18.00 + (\$18.00 / 2) = \$27.00$
- Overtime pay per shift: 2 hours * \$27.00/hour = \$54.00

****Total per day:****

- Regular pay + overtime pay: \$144.00/shift + \$54.00/shift = \$198.00/day

****Total for 5 days:****

- 5 days * \$198.00/day = \$990.00

****Therefore, Tina will make \$990.00 in 5 days.**** The answer is 990.

Q: Abigail is trying a new recipe for a cold drink. It uses 1/4 of a cup of
→ iced tea and 1 and 1/4 of a cup of lemonade to make one drink. If she fills
→ a pitcher with 18 total cups of this drink, how many cups of lemonade are in
→ the pitcher?

A: ## Ambiguity in the Problem Statement:

There is one main ambiguity in the problem statement:

****Total volume vs. Number of servings:**** The statement "18 total cups of this drink" could be interpreted in two ways:

- * ****18 cups of the combined volume:**** This would mean Abigail used a total of 18 cups of liquid, including both iced tea and lemonade.

- * ****18 individual servings:**** This would mean Abigail made 18 individual drinks, each containing $\frac{1}{4}$ cup of iced tea and $1\frac{1}{4}$ cup of lemonade.

Let us assume the interpretation "18 cups of the combined volume".

Solution assuming 18 cups of combined volume:

****Step 1: Find the proportion of lemonade in one drink:****

- * Lemonade: $1\frac{1}{4}$ cups

- * Iced tea: $\frac{1}{4}$ cup

- * Total: $1\frac{1}{4} + \frac{1}{4} = 1\frac{1}{2}$ cups

- * Lemonade proportion: $(1\frac{1}{4}) / (1\frac{1}{2}) = \frac{5}{6}$

****Step 2: Calculate the amount of lemonade in the pitcher:****

- * Total volume: 18 cups

- * Lemonade proportion: $\frac{5}{6}$

- * Volume of lemonade: $18 * (\frac{5}{6}) = 15$ cups

Therefore, there are 15 cups of lemonade in the pitcher. The answer is 15.

Q: A deep-sea monster rises from the waters once every hundred years to feast on a ship and sate its hunger. Over three hundred years, it has consumed 847 people. Ships have been built larger over time, so each new ship has twice as many people as the last ship. How many people were on the ship the monster ate in the first hundred years?

A: Let us solve it using algebra. Let x be the number of people on the ship the monster ate in the first hundred years.

The number of people on the ship eaten in the second hundred years is $2x$, and in the third hundred years is $4x$.

Therefore, the total number of people eaten over three hundred years is $x + 2x + 4x = 847$.

Combining like terms, we get $7x = 847$.

Dividing both sides by 7, we find $x = 121$.

Therefore, there were 121 people on the ship the monster ate in the first hundred years. The answer is 121."

```
[ ]: import re

def extract_number_from_text(text, prefix="The answer is"):
    """
    Extracts the last number from a text string that follows a given prefix.
    Args:
        text (str): The text from which to extract the number.
        prefix (str): The prefix to search for before extracting the number.
    Returns:
        float or None: The extracted number, or None if no valid number is
        found.
    """
    # Find the part of the text that starts with the prefix
    match = re.search(re.escape(prefix) + r".*", text)
    if match:
        # Extract all numbers from the matched text
        numbers = re.findall(r"[-+]?[0-9]*\.[0-9]+", match.group(0))
        if numbers:
            # Return the last number found as a float
            last_number = numbers[-1]
            try:
                return float(last_number)
            except ValueError:
                print(f"Could not convert '{last_number}' to float.")
                return None
    return None

def extract_response_after_question(full_output, question):
    """
    Extracts the line immediately following the question in the model's output.

    Args:
        - full_output (str): The complete output from the model.
        - question (str): The question text used to locate the response line.

    Returns:
        - str: The line following the question line or None if not found.
    """
    # Normalize line breaks
    full_output = full_output.replace('\r\n', '\n').replace('\r', '\n')
    lines = full_output.split('\n')
```

```

# Attempt to find the line containing the question
for i, line in enumerate(lines):
    if question in line:
        # Return the next line if it exists
        if i + 1 < len(lines):
            return lines[i + 1].strip()
        break

return None

# # Example Usage:
# full_output = """
# Q: How many eggs do Janet's ducks lay?
# A: Janet's ducks lay 16 eggs per day. She eats 3 for breakfast.
# She bakes muffins with 4. She sells the rest for $2 per fresh duck egg.
# So, she gets 16 * 3 - 4 * 2 = $48. The answer is $48.
# """
# question = "How many eggs do Janet's ducks lay?"

# next_line = extract_response_after_question(full_output, question)
# print(f"Response after the question: '{next_line}'")

```

```

[ ]: import torch
from datasets import load_dataset

TEMPLATE = """
Q: {}
A: """

# Load GSM8K dataset
gsm8k_test = load_dataset("gsm8k", "main", split="test")

# Assuming model and tokenizer are already initialized
model.eval() # Set the model to evaluation mode

# Helper function to encode inputs
def prepare_input(p):
    # print(question)
    prompt = (PREAMBLE + '\n\n' + PROMPT + '\n' +
              TEMPLATE.format(p))
    return tokenizer(prompt, return_tensors='pt').input_ids

# Function to decode model output
def decode_output(output_ids):
    return tokenizer.decode(output_ids, skip_special_tokens=True)

```

```

# Manual testing loop
all_correct = 0
all_responses = {}
idx = 0
total = len(gsm8k_test)
total = 100

for task_id, problem in enumerate(gsm8k_test):
    if idx == total:
        break

    print(f"task_id {task_id}")

    # Prepare the input for the model
    input_ids = prepare_input(problem['question'])
    with torch.no_grad():
        output_ids = model.generate(input_ids, max_new_tokens = 120) # Adjust
        ↪ max_length as needed

    response = decode_output(output_ids[0])
    all_responses[task_id] = response

    answer_line = extract_response_after_question(response, problem['question'])

    # Compare model output to the ground truth
    model_number = extract_number_from_text(answer_line, "The answer is")
    ground_truth_number = extract_number_from_text(problem['answer'], "####")

    # print(model_number)
    # print(ground_truth_number)

    if model_number == ground_truth_number:
        all_correct += 1

    print(f"Model answer: {model_number}")
    print(f"Ground truth answer: {ground_truth_number}")
    print(f"Correct: {all_correct} out of {total}")
    print("="*40)
    idx += 1

# Final accuracy
accuracy = all_correct / len(gsm8k_test)
print(f"Final Accuracy: {accuracy:.2f}")

```

The attention mask and the pad token id were not set. As a consequence, you may

observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

task_id 0

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 7.0

Ground truth answer: 18.0

Correct: 0 out of 100

=====

task_id 1

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 2.5

Ground truth answer: 3.0

Correct: 0 out of 100

=====

task_id 2

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 20000.0

Ground truth answer: 70000.0

Correct: 0 out of 100

=====

task_id 3

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 180.0

Ground truth answer: 540.0

Correct: 0 out of 100

=====

task_id 4

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain

```

reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 20.0
Correct: 0 out of 100
=====
task_id 5

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 85.0
Ground truth answer: 64.0
Correct: 0 out of 100
=====
task_id 6

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 260.0
Correct: 0 out of 100
=====
task_id 7

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 1.25
Ground truth answer: 160.0
Correct: 0 out of 100
=====
task_id 8

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 3.0
Ground truth answer: 45.0
Correct: 0 out of 100
=====
task_id 9

```


The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 55.0

Ground truth answer: 460.0

Correct: 0 out of 100

=====

task_id 10

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 274.0

Ground truth answer: 366.0

Correct: 0 out of 100

=====

task_id 11

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 360.0

Ground truth answer: 694.0

Correct: 0 out of 100

=====

task_id 12

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 9.0

Ground truth answer: 13.0

Correct: 0 out of 100

=====

task_id 13

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 10.0

Ground truth answer: 18.0

Correct: 0 out of 100

```

=====
task_id 14

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 3.7
Ground truth answer: 60.0
Correct: 0 out of 100
=====
task_id 15

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 125.0
Correct: 0 out of 100
=====
task_id 16

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 150.0
Ground truth answer: 230.0
Correct: 0 out of 100
=====
task_id 17

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 11500.0
Ground truth answer: 57500.0
Correct: 0 out of 100
=====
task_id 18

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

```

Model answer: 12.0

Ground truth answer: 7.0

Correct: 0 out of 100

=====

task_id 19

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None

Ground truth answer: 6.0

Correct: 0 out of 100

=====

task_id 20

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 26.0

Ground truth answer: 15.0

Correct: 0 out of 100

=====

task_id 21

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 8.0

Ground truth answer: 14.0

Correct: 0 out of 100

=====

task_id 22

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 7.0

Ground truth answer: 7.0

Correct: 1 out of 100

=====

task_id 23

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain

reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 8.0
Ground truth answer: 8.0
Correct: 2 out of 100
=====

task_id 24

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 16.25
Ground truth answer: 26.0
Correct: 2 out of 100
=====

task_id 25

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 1.0
Ground truth answer: 2.0
Correct: 2 out of 100
=====

task_id 26

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 156.5
Ground truth answer: 243.0
Correct: 2 out of 100
=====

task_id 27

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 240.0
Ground truth answer: 16.0
Correct: 2 out of 100
=====

task_id 28

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 35.0

Ground truth answer: 25.0

Correct: 2 out of 100

=====

task_id 29

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 66.0

Ground truth answer: 104.0

Correct: 2 out of 100

=====

task_id 30

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 15.0

Ground truth answer: 109.0

Correct: 2 out of 100

=====

task_id 31

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 60.0

Ground truth answer: 80.0

Correct: 2 out of 100

=====

task_id 32

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 120.0

Ground truth answer: 35.0

Correct: 2 out of 100

```

=====
task_id 33

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 30.0
Ground truth answer: 70.0
Correct: 2 out of 100
=====
task_id 34

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 38.0
Ground truth answer: 23.0
Correct: 2 out of 100
=====
task_id 35

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 5.25
Ground truth answer: 9.0
Correct: 2 out of 100
=====
task_id 36

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 120.0
Ground truth answer: 75.0
Correct: 2 out of 100
=====
task_id 37

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

```

Model answer: 5.0

Ground truth answer: 2.0

Correct: 2 out of 100

=====

task_id 38

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None

Ground truth answer: 10.0

Correct: 2 out of 100

=====

task_id 39

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None

Ground truth answer: 18.0

Correct: 2 out of 100

=====

task_id 40

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 4.0

Ground truth answer: 8.0

Correct: 2 out of 100

=====

task_id 41

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 1200.0

Ground truth answer: 200.0

Correct: 2 out of 100

=====

task_id 42

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain

reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 40.0
Ground truth answer: 26.0
Correct: 2 out of 100
=====

task_id 43

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 48.0
Correct: 2 out of 100
=====

task_id 44

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 20.0
Correct: 2 out of 100
=====

task_id 45

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 104.0
Correct: 2 out of 100
=====

task_id 46

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 220.0
Ground truth answer: 163.0
Correct: 2 out of 100
=====

task_id 47

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None

Ground truth answer: 800.0

Correct: 2 out of 100

=====

task_id 48

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 7.0

Ground truth answer: 8.0

Correct: 2 out of 100

=====

task_id 49

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 108.0

Ground truth answer: 30.0

Correct: 2 out of 100

=====

task_id 50

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 307.5

Ground truth answer: 294.0

Correct: 2 out of 100

=====

task_id 51

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 1.5

Ground truth answer: 5.0

Correct: 2 out of 100

```

=====
task_id 52

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 15.0
Correct: 2 out of 100
=====
task_id 53

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 40.0
Correct: 2 out of 100
=====
task_id 54

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 40.0
Correct: 2 out of 100
=====
task_id 55

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 7.0
Ground truth answer: 14.0
Correct: 2 out of 100
=====
task_id 56

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

```

Model answer: 3.0

Ground truth answer: 3.0

Correct: 3 out of 100

=====

task_id 57

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 30.0

Ground truth answer: 83.0

Correct: 3 out of 100

=====

task_id 58

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 40.0

Ground truth answer: 57.0

Correct: 3 out of 100

=====

task_id 59

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 133.0

Ground truth answer: 187.0

Correct: 3 out of 100

=====

task_id 60

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 18.0

Ground truth answer: 17.0

Correct: 3 out of 100

=====

task_id 61

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain

reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 1000.0
Ground truth answer: 1430.0
Correct: 3 out of 100
=====

task_id 62

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 25000.0
Correct: 3 out of 100
=====

task_id 63

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 1596.0
Correct: 3 out of 100
=====

task_id 64

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 25.0
Ground truth answer: 300.0
Correct: 3 out of 100
=====

task_id 65

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 30.0
Ground truth answer: 36.0
Correct: 3 out of 100
=====

task_id 66

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 80.0

Ground truth answer: 48.0

Correct: 3 out of 100

=====

task_id 67

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None

Ground truth answer: 595.0

Correct: 3 out of 100

=====

task_id 68

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 12.0

Ground truth answer: 36.0

Correct: 3 out of 100

=====

task_id 69

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 95.0

Ground truth answer: 60.0

Correct: 3 out of 100

=====

task_id 70

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 50.0

Ground truth answer: 7425.0

Correct: 3 out of 100

```

=====
task_id 71

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 60.0
Ground truth answer: 60.0
Correct: 4 out of 100
=====
task_id 72

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 105.0
Ground truth answer: 221.0
Correct: 4 out of 100
=====
task_id 73

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 1.5
Ground truth answer: 255.0
Correct: 4 out of 100
=====
task_id 74

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 88.0
Correct: 4 out of 100
=====
task_id 75

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

```

Model answer: 4.0

Ground truth answer: 60.0

Correct: 4 out of 100

=====

task_id 76

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None

Ground truth answer: 5.0

Correct: 4 out of 100

=====

task_id 77

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 400.0

Ground truth answer: 100.0

Correct: 4 out of 100

=====

task_id 78

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 36.0

Ground truth answer: 6.0

Correct: 4 out of 100

=====

task_id 79

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 70.0

Ground truth answer: 70.0

Correct: 5 out of 100

=====

task_id 80

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain

reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 14.0
Ground truth answer: 10.0
Correct: 5 out of 100
=====

task_id 81

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 1.0
Ground truth answer: 17.0
Correct: 5 out of 100
=====

task_id 82

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 92.0
Ground truth answer: 623.0
Correct: 5 out of 100
=====

task_id 83

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 375.0
Ground truth answer: 600.0
Correct: 5 out of 100
=====

task_id 84

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 14.0
Ground truth answer: 15.0
Correct: 5 out of 100
=====

task_id 85

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None

Ground truth answer: 44.0

Correct: 5 out of 100

=====

task_id 86

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 8.0

Ground truth answer: 22.0

Correct: 5 out of 100

=====

task_id 87

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 720.0

Ground truth answer: 9360.0

Correct: 5 out of 100

=====

task_id 88

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 800.0

Ground truth answer: 8000.0

Correct: 5 out of 100

=====

task_id 89

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.

Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 24.0

Ground truth answer: 24.0

Correct: 6 out of 100

```

=====
task_id 90

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 210.0
Ground truth answer: 225.0
Correct: 6 out of 100
=====
task_id 91

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 6.0
Ground truth answer: 28.0
Correct: 6 out of 100
=====
task_id 92

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: -1.0
Ground truth answer: 4.0
Correct: 6 out of 100
=====
task_id 93

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: None
Ground truth answer: 36.0
Correct: 6 out of 100
=====
task_id 94

The attention mask and the pad token id were not set. As a consequence, you may
observe unexpected behavior. Please pass your input's `attention_mask` to obtain
reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

```

Model answer: 108.0
Ground truth answer: 348.0
Correct: 6 out of 100
=====

task_id 95

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 198.0
Ground truth answer: 40.0
Correct: 6 out of 100
=====

task_id 96

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 3.0
Ground truth answer: 3.0
Correct: 7 out of 100
=====

task_id 97

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 2.0
Ground truth answer: 12.0
Correct: 7 out of 100
=====

task_id 98

The attention mask and the pad token id were not set. As a consequence, you may observe unexpected behavior. Please pass your input's `attention_mask` to obtain reliable results.
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.

Model answer: 35.0
Ground truth answer: 5.0
Correct: 7 out of 100
=====

task_id 99

Model answer: 0.0
Ground truth answer: 58.0
Correct: 7 out of 100

=====

Final Accuracy: 0.01

```
[ ]: accuracy = all_correct / total* 100
      print(f"Final Accuracy codellama pretrained on trace : {accuracy:.2f}%")
```

Final Accuracy codellama pretrained on trace : 7.00%

```
[ ]:
```