

llama_trace

April 12, 2024

```
[ ]: # !pip install -q -U bitsandbytes
# !pip install -q -U git+https://github.com/huggingface/transformers.git
# !pip install -q -U git+https://github.com/huggingface/peft.git
# !pip install -q -U git+https://github.com/huggingface/accelerate.git
# !pip install -q -U datasets scipy ipywidgets matplotlib

[ ]: from unsloth import FastLanguageModel
import torch
max_seq_length = 2048 # Choose any! We auto support RoPE Scaling internally!
dtype = None # None for auto detection. Float16 for Tesla T4, V100, Bfloat16
    ↪ for Ampere+
load_in_4bit = True # Use 4bit quantization to reduce memory usage. Can be
    ↪ False.

# 4bit pre quantized models we support for 4x faster downloading + no OOMs.
fourbit_models = [
    "unsloth/mistral-7b-bnb-4bit",
    "unsloth/mistral-7b-instruct-v0.2-bnb-4bit",
    "unsloth/llama-2-7b-bnb-4bit",
    "unsloth/llama-2-13b-bnb-4bit",
    "unsloth/codellama-34b-bnb-4bit",
    "unsloth/tinyllama-bnb-4bit",
    "unsloth/gemma-7b-bnb-4bit", # New Google 6 trillion tokens model 2.5x
    ↪ faster!
    "unsloth/gemma-2b-bnb-4bit",
] # More models at https://huggingface.co/unsloth

# unsloth/codellama-7b-bnb-4bit

model, tokenizer = FastLanguageModel.from_pretrained(
    model_name = "unsloth/llama-2-7b-bnb-4bit", # Choose ANY! eg mistralai/
    ↪ Mistral-7B-Instruct-v0.2
    max_seq_length = max_seq_length,
    dtype = dtype,
    load_in_4bit = load_in_4bit,
    # token = "hf_...", # use one if using gated models like meta-llama/
    ↪ Llama-2-7b-hf
```

```
)
```

```
config.json: 0%|          | 0.00/1.10k [00:00<?, ?B/s]
Unused kwargs: ['quant_method']. These kwargs are not used in <class
'transformers.utils.quantization_config.BitsAndBytesConfig'>.

==(====)== Unsloth: Fast Llama patching release 2024.3
  \ \ / | GPU: NVIDIA GeForce RTX 3090. Max memory: 23.483 GB. Platform =
Linux.
0^0/ \_/ \ Pytorch: 2.2.1+cu121. CUDA = 8.6. CUDA Toolkit = 12.1.
\      / Bfloat16 = TRUE. Xformers = 0.0.25. FA = True.
"-____-" Free Apache license: http://github.com/unslothai/unsloth

model.safetensors: 0%|          | 0.00/3.87G [00:00<?, ?B/s]
generation_config.json: 0%|          | 0.00/188 [00:00<?, ?B/s]
tokenizer_config.json: 0%|          | 0.00/894 [00:00<?, ?B/s]
tokenizer.model: 0%|          | 0.00/500k [00:00<?, ?B/s]
tokenizer.json: 0%|          | 0.00/1.84M [00:00<?, ?B/s]
special_tokens_map.json: 0%|          | 0.00/438 [00:00<?, ?B/s]
```

```
[ ]: model = FastLanguageModel.get_peft_model(
    model,
    r = 16, # Choose any number > 0 ! Suggested 8, 16, 32, 64, 128
    target_modules = ["q_proj", "k_proj", "v_proj", "o_proj",
                      "gate_proj", "up_proj", "down_proj",],
    lora_alpha = 16,
    lora_dropout = 0, # Supports any, but = 0 is optimized
    bias = "none", # Supports any, but = "none" is optimized
    use_gradient_checkpointing = True,
    random_state = 3407,
    use_rslora = False, # We support rank stabilized LoRA
    loftq_config = None, # And LoftQ
)
```

Unsloth 2024.3 patched 32 layers with 32 QKV layers, 32 O layers and 32 MLP layers.

```
[ ]: from datasets import load_dataset
tokenizer.pad_token = tokenizer.eos_token

json_file_path = "./python_states_singleline.json"
alpaca_prompt = """"Below is an instruction that describes a task, paired with
    ↪an input that provides further context. Write a response that appropriately
    ↪completes the request.
```

```

### Instruction:
Write down all of the state changes that take place after the code snippet is
↳executed.

### Input:
{}

### Response:
{}"""
# trace_prompt = """<s>[INST] {} [/INST] {}</s>"""

def formatting_prompts_func(examples):
    inputs = examples["input"]
    outputs = examples["output"]
    texts = []
    for input, output in zip(inputs, outputs):
        text = alpaca_prompt.format(input, output)
        texts.append(text)
    return {"text": texts}

dataset = load_dataset("json", data_files=json_file_path, split="train").
↳select(range(2001))
dataset = dataset.map(formatting_prompts_func, batched=True) # had to unset
↳batched

```

```

[ ]: from trl import SFTTrainer
from transformers import TrainingArguments

trainer = SFTTrainer(
    model = model,
    tokenizer = tokenizer,
    train_dataset = dataset,
    dataset_text_field = "text",
    max_seq_length = max_seq_length,
    dataset_num_proc = 2,
    packing = False, # Can make training 5x faster for short sequences.
    args = TrainingArguments(
        per_device_train_batch_size = 2,
        gradient_accumulation_steps = 4,
        warmup_steps = 5,
        # max_steps = 200,
        num_train_epochs=1,
        learning_rate = 2e-4,
        fp16 = not torch.cuda.is_bf16_supported(),
        bf16 = torch.cuda.is_bf16_supported(),
        logging_steps = 1,

```

```

    optim = "adamw_8bit",
    weight_decay = 0.01,
    lr_scheduler_type = "linear",
    seed = 3407,
    output_dir = "outputs",
),
)

```

Map (num_proc=2): 0% | 0/2001 [00:00<?, ? examples/s]

max_steps is given, it will override any value given in num_train_epochs

```
[ ]: trainer_stats = trainer.train()
```

```

==(====)== Unsloth - 2x faster free finetuning | Num GPUs = 1
  \ \ / | Num examples = 2,001 | Num Epochs = 1
0^0/ \_/ \ Batch size per device = 2 | Gradient Accumulation steps = 4
\         / Total batch size = 8 | Total steps = 200
"-____-" Number of trainable parameters = 39,976,960

<IPython.core.display.HTML object>

```

```

[ ]: obj = {
    "input": "state: h = [None, {}, {}, None, None, None]; j = 2; o = 'k';\n
    ↪code: h[j][o] = 1",
    "output": "h = [None, {}, {'k': 1}, None, None, None]; j = 2; o = 'k';",
    "example": 8585426
}

FastLanguageModel.for_inference(model) # Enable native 2x faster inference
inputs = tokenizer([
    alpaca_prompt.format(
        "state: h = [None, {}, {}, None, None, None]; j = 2; o = 'k'; code:\n
    ↪h[j][o] = 1"
        , "")
], return_tensors = "pt").to('cuda')

```

```

[ ]: outputs = model.generate(**inputs, max_new_tokens = 64, use_cache = True)
print(tokenizer.batch_decode(outputs))

```

["<s> Below is an instruction that describes a task, paired with an input that provides further context. Write a response that appropriately completes the request.\n\n### Instruction:\nWrite down all of the state changes that take place after the code snippet is executed.\n\n### Input:\nstate: h = [None, {}, {}, None, None, None]; j = 2; o = 'k'; code: h[j][o] = 1\n\n### Response:\nh = [None, {}, {}, None, None, 1]; j = 2; o = 'k';\n\n### Explanation:\nh[j][o] = 1;\n\n### Explanation:\nh = [None, {}, {}, None, None, "]

```
[ ]: from collections import Counter

# EVAL LOGIC
def calculate_token_level_f1(prediction_tokens, reference_tokens):
    """
    Calculate precision, recall, and F1 score based on token overlap.
    """
    common_token_count = Counter(prediction_tokens) & Counter(reference_tokens)
    num_same = sum(common_token_count.values())

    if num_same == 0:
        return 0, 0, 0

    precision = 1.0 * num_same / len(prediction_tokens)
    recall = 1.0 * num_same / len(reference_tokens)
    f1 = (2 * precision * recall) / (precision + recall)

    return precision, recall, f1

def correct_solution(prediction_str, reference_str):
    """
    Compare the final numerical output of the model with the reference tokens.

    Args:
    - prediction_tokens: List of token IDs representing the model's prediction.
    - reference_tokens: List of token IDs representing the reference output.

    Returns:
    - 1 if the final numerical output of the model matches the reference tokens
    exactly, else 0.
    """
    # prediction_str = tokenizer.decode(prediction_tokens,
    ↪ skip_special_tokens=True)
    # reference_str = tokenizer.decode(reference_tokens,
    ↪ skip_special_tokens=True)

    prediction_lines = prediction_str.strip().split("\n")
    reference_lines = reference_str.strip().split("\n")

    # print(prediction_lines)
    # print(reference_lines)

    last_prediction_line = prediction_lines[-1].strip()
    last_reference_line = reference_lines[-1].strip()

    # print("predicted ", last_prediction_line)
```

```

# print("reference ",last_reference_line)
# print(last_prediction_line== last_reference_line)

if last_prediction_line== last_reference_line:
    return 1
else:
    return 0

def custom_metrics_gsm8k(preds):
    # TODO Changed this function group to work with gsm8k
    logits = torch.tensor(preds.predictions)
    labels = torch.tensor(preds.label_ids)

    batch_size, seq_length, vocab_size = logits.shape

    # steal from inside llama
    # shift logits by 1 index cuz of causal lm
    shift_logits = logits[..., :-1, :].contiguous()
    shift_labels = labels[..., 1:].contiguous()
    # Flatten the tokens
    # loss_fct = CrossEntropyLoss()
    shift_logits = shift_logits.view(batch_size, -1, vocab_size)
    shift_labels = shift_labels.view(batch_size, -1)

    probs = torch.nn.functional.softmax(shift_logits.view(-1, vocab_size),
    ↪dim=-1)
    p_true_tokens = probs.view(-1, vocab_size)[
        torch.arange(batch_size * (seq_length - 1)), shift_labels.view(-1)
    ].view(batch_size, (seq_length - 1))

    nll = -torch.log(p_true_tokens)
    mean_nll = nll.mean()
    ppl = torch.exp(mean_nll) # perplexity

    # compute percentage of correct tokens
    correct_tokens = (
        (shift_logits.view(-1, vocab_size).argmax(-1) == shift_labels.view(-1))
        .float()
        .mean()
    )

    pred_max_labels = shift_logits.argmax(-1).view(batch_size, -1)

    f1_scores = []
    precision_scores = []
    recall_scores = []

```

```

solution_scores = []

for i in range(batch_size):
    unmasked_label_tokens = shift_labels[i][shift_labels[i] != -100][
        :-1
    ] # drop eos_token
    # find the index where the instruction token ends and the answer begins
    inst_token_seq = tokenizer.encode("/INST", return_tensors="pt")[0][1:]
    first_output_idx = None
    for j in range(unmasked_label_tokens.shape[0] - len(inst_token_seq)):
        if torch.equal(
            unmasked_label_tokens[j : j + len(inst_token_seq)],
            inst_token_seq
        ):
            first_output_idx = j + len(inst_token_seq)
            break
    assert (
        first_output_idx is not None
    ), "Could not find the end of the instruction token"

    # get ground truth output tokens
    gt_output_tokens = unmasked_label_tokens[first_output_idx:]
    # get predicted output tokens (including padding)
    pred_output_tokens_masked = pred_max_labels[i][first_output_idx:]
    # drop the pad tokens
    pred_output_tokens_unmasked = pred_output_tokens_masked[
        pred_output_tokens_masked != -100
    ]

    eos_token_indices = torch.where(
        pred_output_tokens_unmasked == tokenizer.eos_token_id
    )[0]

    if eos_token_indices.size(0) > 0:
        first_pred_output_stop_idx = eos_token_indices[0].item()
    else:
        first_pred_output_stop_idx = len(pred_output_tokens_unmasked) - 1

    pred_output_tokens = pred_output_tokens_unmasked[:
        first_pred_output_stop_idx]

    gt_output_str = tokenizer.decode(gt_output_tokens)
    pred_output_str = tokenizer.decode(pred_output_tokens)

    precision, recall, f1 = calculate_token_level_f1(pred_output_str,
        gt_output_str)

```

```

        correct = correct_solution(pred_output_str, gt_output_str)
        solution_scores.append(correct)

    f1_scores.append(f1)
    precision_scores.append(precision)
    recall_scores.append(recall)

mean_f1 = np.mean(f1_scores) if f1_scores else 0
mean_precision = np.mean(precision_scores) if precision_scores else 0
mean_recall = np.mean(recall_scores) if recall_scores else 0
solve_rate = np.mean(solution_scores) if solution_scores else 0

# wandb.log(
#     {
#         "perplexity": ppl.item(),
#         "correct_tokens": correct_tokens.item(),
#         "f1": mean_f1,
#         "solve_rate": solve_rate,
#     }
# )
return {
    "perplexity": ppl,
    "correct_tokens": correct_tokens.item(),
    "f1": mean_f1,
    "mean_precision": mean_precision,
    "mean_recall": mean_recall,
    "solve_rate": solve_rate,
}

```

```
[ ]: # @title GSM8K Prompts
```

```

PREAMBLE = """As an expert problem solver solve step by step the following
↳mathematical questions."""

```

```

# The default gsm8k prompt from the CoT paper
# https://arxiv.org/pdf/2201.11903.pdf page 35.

```

```

PROMPT = """Q: There are 15 trees in the grove. Grove workers will plant trees
↳in the grove today. After they are done, there will be 21 trees. How many
↳trees did the grove workers plant today?

```

```

A: We start with 15 trees. Later we have 21 trees. The difference must be the
↳number of trees they planted. So, they must have planted  $21 - 15 = 6$  trees.
↳The answer is 6.

```

```

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars
↳are in the parking lot?

```


A: There are 3 cars in the parking lot already. 2 more arrive. Now there are 3
→ $3 + 2 = 5$ cars. The answer is 5.

Q: Leah had 32 chocolates and her sister had 42. If they ate 35, how many
→ pieces do they have left in total?

A: Leah had 32 chocolates and Leah's sister had 42. That means there were
→ originally $32 + 42 = 74$ chocolates. 35 have been eaten. So in total they
→ still have $74 - 35 = 39$ chocolates. The answer is 39.

Q: Jason had 20 lollipops. He gave Denny some lollipops. Now Jason has 12
→ lollipops. How many lollipops did Jason give to Denny?

A: Jason had 20 lollipops. Since he only has 12 now, he must have given the
→ rest to Denny. The number of lollipops he has given to Denny must have been
→ $20 - 12 = 8$ lollipops. The answer is 8.

Q: Shawn has five toys. For Christmas, he got two toys each from his mom and
→ dad. How many toys does he have now?

A: He has 5 toys. He got 2 from mom, so after that he has $5 + 2 = 7$ toys. Then
→ he got 2 more from dad, so in total he has $7 + 2 = 9$ toys. The answer is 9.

Q: There were nine computers in the server room. Five more computers were
→ installed each day, from monday to thursday. How many computers are now in
→ the server room?

A: There are 4 days from monday to thursday. 5 computers were added each day.
→ That means in total $4 * 5 = 20$ computers were added. There were 9 computers
→ in the beginning, so now there are $9 + 20 = 29$ computers. The answer is 29.

Q: Michael had 58 golf balls. On tuesday, he lost 23 golf balls. On wednesday,
→ he lost 2 more. How many golf balls did he have at the end of wednesday?

A: Michael initially had 58 balls. He lost 23 on Tuesday, so after that he has
→ $58 - 23 = 35$ balls. On Wednesday he lost 2 more so now he has $35 - 2 = 33$
→ balls. The answer is 33.

Q: Olivia has \$23. She bought five bagels for \$3 each. How much money does she
→ have left?

A: She bought 5 bagels for \$3 each. This means she spent $5 * \$3 = \15 on the
→ bagels. She had \$23 in beginning, so now she has $\$23 - \$15 = \$8$. The answer
→ is 8. ""

*# Extension of the default 8-shot prompt, page 35 in
<https://arxiv.org/pdf/2201.11903.pdf>
The extension is intended to improve performance on
more complicated gsm8k examples.*

EXTRA_3_SHOTS = ""As an expert problem solver solve step by step the following
↳mathematical questions.

Q: Tina makes \$18.00 an hour. If she works more than 8 hours per shift, she is
↳eligible for overtime, which is paid by your hourly wage + 1/2 your hourly
↳wage. If she works 10 hours every day for 5 days, how much money does she
↳make?

A: Here's how to calculate Tina's earnings:

****Regular Time:****

- Hours per shift: 8 hours
- Wage per hour: \$18.00
- Regular pay per shift: 8 hours * \$18.00/hour = \$144.00

****Overtime:****

- Overtime hours per shift: 10 hours - 8 hours = 2 hours
- Overtime pay per hour: \$18.00 + (\$18.00 / 2) = \$27.00
- Overtime pay per shift: 2 hours * \$27.00/hour = \$54.00

****Total per day:****

- Regular pay + overtime pay: \$144.00/shift + \$54.00/shift = \$198.00/day

****Total for 5 days:****

- 5 days * \$198.00/day = \$990.00

****Therefore, Tina will make \$990.00 in 5 days.**** The answer is 990.

Q: Abigail is trying a new recipe for a cold drink. It uses 1/4 of a cup of
↳iced tea and 1 and 1/4 of a cup of lemonade to make one drink. If she fills
↳a pitcher with 18 total cups of this drink, how many cups of lemonade are in
↳the pitcher?

A: ## Ambiguity in the Problem Statement:

There is one main ambiguity in the problem statement:

****Total volume vs. Number of servings:**** The statement "18 total cups of this
↳drink" could be interpreted in two ways:

* ****18 cups of the combined volume:**** This would mean Abigail used a total of
↳18 cups of liquid, including both iced tea and lemonade.

* ****18 individual servings:**** This would mean Abigail made 18 individual
↳drinks, each containing 1/4 cup of iced tea and 1 1/4 cup of lemonade.

Let us assume the interpretation "18 cups of the combined volume".

Solution assuming 18 cups of combined volume:

****Step 1: Find the proportion of lemonade in one drink:****

- * Lemonade: $1 \frac{1}{4}$ cups
- * Iced tea: $\frac{1}{4}$ cup
- * Total: $1 \frac{1}{4} + \frac{1}{4} = 1 \frac{1}{2}$ cups
- * Lemonade proportion: $(1 \frac{1}{4}) / (1 \frac{1}{2}) = \frac{5}{6}$

****Step 2: Calculate the amount of lemonade in the pitcher:****

- * Total volume: 18 cups
- * Lemonade proportion: $\frac{5}{6}$
- * Volume of lemonade: $18 * (\frac{5}{6}) = 15$ cups

Therefore, there are 15 cups of lemonade in the pitcher. The answer is 15.

Q: A deep-sea monster rises from the waters once every hundred years to feast
↳ on a ship and sate its hunger. Over three hundred years, it has consumed 847
↳ people. Ships have been built larger over time, so each new ship has twice
↳ as many people as the last ship. How many people were on the ship the
↳ monster ate in the first hundred years?

A: Let us solve it using algebra. Let x be the number of people on the ship the
↳ monster ate in the first hundred years.

The number of people on the ship eaten in the second hundred years is $2x$, and
↳ in the third hundred years is $4x$.

Therefore, the total number of people eaten over three hundred years is $x + 2x$
↳ $+ 4x = 847$.

Combining like terms, we get $7x = 847$.

Dividing both sides by 7, we find $x = 121$.

Therefore, there were 121 people on the ship the monster ate in the first
↳ hundred years. The answer is 121."

```
[ ]: import re

def extract_number_from_text(text, prefix="The answer is"):
    """
    Extracts the last number from a text string that follows a given prefix.
    Args:
        text (str): The text from which to extract the number.
        prefix (str): The prefix to search for before extracting the number.
    Returns:
        float or None: The extracted number, or None if no valid number is
        ↳ found.
```

```

"""
# Find the part of the text that starts with the prefix
match = re.search(re.escape(prefix) + r".*", text)
if match:
    # Extract all numbers from the matched text
    numbers = re.findall(r"[-+]?[0-9]*\.[0-9]+", match.group(0))
    if numbers:
        # Return the last number found as a float
        last_number = numbers[-1]
        try:
            return float(last_number)
        except ValueError:
            print(f"Could not convert '{last_number}' to float.")
            return None
    return None

def extract_response_after_question(full_output, question):
    """
    Extracts the line immediately following the question in the model's output.

    Args:
    - full_output (str): The complete output from the model.
    - question (str): The question text used to locate the response line.

    Returns:
    - str: The line following the question line or None if not found.
    """
    # Normalize line breaks
    full_output = full_output.replace('\r\n', '\n').replace('\r', '\n')
    lines = full_output.split('\n')

    # Attempt to find the line containing the question
    for i, line in enumerate(lines):
        if question in line:
            # Return the next line if it exists
            if i + 1 < len(lines):
                return lines[i + 1].strip()
            break

    return None

# # Example Usage:
# full_output = """
# Q: How many eggs do Janet's ducks lay?
# A: Janet's ducks lay 16 eggs per day. She eats 3 for breakfast.
# She bakes muffins with 4. She sells the rest for $2 per fresh duck egg.
# So, she gets 16 * 3 - 4 * 2 = $48. The answer is $48.

```

```

# """
# question = "How many eggs do Janet's ducks lay?"

# next_line = extract_response_after_question(full_output, question)
# print(f"Response after the question: '{next_line}'")

```

```

[ ]: import torch
from datasets import load_dataset

TEMPLATE = """
Q: {}
A: """

# Load GSM8K dataset
gsm8k_test = load_dataset("gsm8k", "main", split="test")

# Assuming model and tokenizer are already initialized
model.eval() # Set the model to evaluation mode

# Helper function to encode inputs
def prepare_input(p):
    # print(question)
    prompt = (PREAMBLE + '\n\n' + PROMPT + '\n' +
              TEMPLATE.format(p))
    return tokenizer(prompt, return_tensors='pt').input_ids

# Function to decode model output
def decode_output(output_ids):
    return tokenizer.decode(output_ids, skip_special_tokens=True)

# Manual testing loop
all_correct = 0
all_responses = {}
idx = 0
total = len(gsm8k_test)
total = 100

for task_id, problem in enumerate(gsm8k_test):
    if idx == total:
        break

    print(f"task_id {task_id}")

    # Prepare the input for the model
    input_ids = prepare_input(problem['question'])
    with torch.no_grad():

```

```

        output_ids = model.generate(input_ids, max_new_tokens = 120) # Adjust
        ↪max_length as needed

    response = decode_output(output_ids[0])
    all_responses[task_id] = response

    answer_line = extract_response_after_question(response, problem['question'])

    # Compare model output to the ground truth
    model_number = extract_number_from_text(answer_line, "The answer is")
    ground_truth_number = extract_number_from_text(problem['answer'], "####")

    # print(model_number)
    # print(ground_truth_number)

    if model_number == ground_truth_number:
        all_correct += 1

    print(f"Model answer: {model_number}")
    print(f"Ground truth answer: {ground_truth_number}")
    print(f"Correct: {all_correct} out of {total}")
    print("="*40)
    idx += 1

# Final accuracy
accuracy = all_correct / len(gsm8k_test)
print(f"Final Accuracy: {accuracy:.2f}")

```

```

task_id 0
Model answer: None
Ground truth answer: 18.0
Correct: 0 out of 100
=====
task_id 1
Model answer: 0.0
Ground truth answer: 3.0
Correct: 0 out of 100
=====
task_id 2
Model answer: 0.0
Ground truth answer: 70000.0
Correct: 0 out of 100
=====
task_id 3
Model answer: 540.0

```

```

Ground truth answer: 540.0
Correct: 1 out of 100
=====
task_id 4
Model answer: 240.0
Ground truth answer: 20.0
Correct: 1 out of 100
=====
task_id 5
Model answer: 80.0
Ground truth answer: 64.0
Correct: 1 out of 100
=====
task_id 6
Model answer: None
Ground truth answer: 260.0
Correct: 1 out of 100
=====
task_id 7
Model answer: 0.0
Ground truth answer: 160.0
Correct: 1 out of 100
=====
task_id 8
Model answer: None
Ground truth answer: 45.0
Correct: 1 out of 100
=====
task_id 9
Model answer: None
Ground truth answer: 460.0
Correct: 1 out of 100
=====
task_id 10
Model answer: None
Ground truth answer: 366.0
Correct: 1 out of 100
=====
task_id 11
Model answer: None
Ground truth answer: 694.0
Correct: 1 out of 100
=====
task_id 12
Model answer: None
Ground truth answer: 13.0
Correct: 1 out of 100
=====

```

```

task_id 13
Model answer: 5.0
Ground truth answer: 18.0
Correct: 1 out of 100
=====
task_id 14
Model answer: None
Ground truth answer: 60.0
Correct: 1 out of 100
=====
task_id 15
Model answer: None
Ground truth answer: 125.0
Correct: 1 out of 100
=====
task_id 16
Model answer: 7.0
Ground truth answer: 230.0
Correct: 1 out of 100
=====
task_id 17
Model answer: None
Ground truth answer: 57500.0
Correct: 1 out of 100
=====
task_id 18
Model answer: 24.0
Ground truth answer: 7.0
Correct: 1 out of 100
=====
task_id 19
Model answer: 3.0
Ground truth answer: 6.0
Correct: 1 out of 100
=====
task_id 20
Model answer: None
Ground truth answer: 15.0
Correct: 1 out of 100
=====
task_id 21
Model answer: 29.0
Ground truth answer: 14.0
Correct: 1 out of 100
=====
task_id 22
Model answer: 7.0
Ground truth answer: 7.0

```



```

Correct: 2 out of 100
=====
task_id 23
Model answer: None
Ground truth answer: 8.0
Correct: 2 out of 100
=====
task_id 24
Model answer: 24.38
Ground truth answer: 26.0
Correct: 2 out of 100
=====
task_id 25
Model answer: None
Ground truth answer: 2.0
Correct: 2 out of 100
=====
task_id 26
Model answer: None
Ground truth answer: 243.0
Correct: 2 out of 100
=====
task_id 27
Model answer: None
Ground truth answer: 16.0
Correct: 2 out of 100
=====
task_id 28
Model answer: 25.0
Ground truth answer: 25.0
Correct: 3 out of 100
=====
task_id 29
Model answer: 96.0
Ground truth answer: 104.0
Correct: 3 out of 100
=====
task_id 30
Model answer: 174.0
Ground truth answer: 109.0
Correct: 3 out of 100
=====
task_id 31
Model answer: None
Ground truth answer: 80.0
Correct: 3 out of 100
=====
task_id 32

```

```
Model answer: 35.0
Ground truth answer: 35.0
Correct: 4 out of 100
=====
task_id 33
Model answer: 140.0
Ground truth answer: 70.0
Correct: 4 out of 100
=====
task_id 34
Model answer: None
Ground truth answer: 23.0
Correct: 4 out of 100
=====
task_id 35
Model answer: 5.0
Ground truth answer: 9.0
Correct: 4 out of 100
=====
task_id 36
Model answer: 120.0
Ground truth answer: 75.0
Correct: 4 out of 100
=====
task_id 37
Model answer: None
Ground truth answer: 2.0
Correct: 4 out of 100
=====
task_id 38
Model answer: None
Ground truth answer: 10.0
Correct: 4 out of 100
=====
task_id 39
Model answer: None
Ground truth answer: 18.0
Correct: 4 out of 100
=====
task_id 40
Model answer: None
Ground truth answer: 8.0
Correct: 4 out of 100
=====
task_id 41
Model answer: None
Ground truth answer: 200.0
Correct: 4 out of 100
```

```

=====
task_id 42
Model answer: 9.0
Ground truth answer: 26.0
Correct: 4 out of 100
=====
task_id 43
Model answer: None
Ground truth answer: 48.0
Correct: 4 out of 100
=====
task_id 44
Model answer: None
Ground truth answer: 20.0
Correct: 4 out of 100
=====
task_id 45
Model answer: None
Ground truth answer: 104.0
Correct: 4 out of 100
=====
task_id 46
Model answer: None
Ground truth answer: 163.0
Correct: 4 out of 100
=====
task_id 47
Model answer: 5.0
Ground truth answer: 800.0
Correct: 4 out of 100
=====
task_id 48
Model answer: 24.0
Ground truth answer: 8.0
Correct: 4 out of 100
=====
task_id 49
Model answer: 120.0
Ground truth answer: 30.0
Correct: 4 out of 100
=====
task_id 50
Model answer: None
Ground truth answer: 294.0
Correct: 4 out of 100
=====
task_id 51
Model answer: 46.0

```

```

Ground truth answer: 5.0
Correct: 4 out of 100
=====
task_id 52
Model answer: 7.0
Ground truth answer: 15.0
Correct: 4 out of 100
=====
task_id 53
Model answer: 120.0
Ground truth answer: 40.0
Correct: 4 out of 100
=====
task_id 54
Model answer: 33.0
Ground truth answer: 40.0
Correct: 4 out of 100
=====
task_id 55
Model answer: 56.0
Ground truth answer: 14.0
Correct: 4 out of 100
=====
task_id 56
Model answer: 2.0
Ground truth answer: 3.0
Correct: 4 out of 100
=====
task_id 57
Model answer: None
Ground truth answer: 83.0
Correct: 4 out of 100
=====
task_id 58
Model answer: 54.0
Ground truth answer: 57.0
Correct: 4 out of 100
=====
task_id 59
Model answer: 124.0
Ground truth answer: 187.0
Correct: 4 out of 100
=====
task_id 60
Model answer: 20.0
Ground truth answer: 17.0
Correct: 4 out of 100
=====

```

```

task_id 61
Model answer: 930.0
Ground truth answer: 1430.0
Correct: 4 out of 100
=====
task_id 62
Model answer: None
Ground truth answer: 25000.0
Correct: 4 out of 100
=====
task_id 63
Model answer: 266.0
Ground truth answer: 1596.0
Correct: 4 out of 100
=====
task_id 64
Model answer: None
Ground truth answer: 300.0
Correct: 4 out of 100
=====
task_id 65
Model answer: 48.0
Ground truth answer: 36.0
Correct: 4 out of 100
=====
task_id 66
Model answer: None
Ground truth answer: 48.0
Correct: 4 out of 100
=====
task_id 67
Model answer: 605.0
Ground truth answer: 595.0
Correct: 4 out of 100
=====
task_id 68
Model answer: 900.0
Ground truth answer: 36.0
Correct: 4 out of 100
=====
task_id 69
Model answer: 60.0
Ground truth answer: 60.0
Correct: 5 out of 100
=====
task_id 70
Model answer: None
Ground truth answer: 7425.0

```

```
Correct: 5 out of 100
=====
task_id 71
Model answer: 40.0
Ground truth answer: 60.0
Correct: 5 out of 100
=====
task_id 72
Model answer: 234.0
Ground truth answer: 221.0
Correct: 5 out of 100
=====
task_id 73
Model answer: 45.0
Ground truth answer: 255.0
Correct: 5 out of 100
=====
task_id 74
Model answer: None
Ground truth answer: 88.0
Correct: 5 out of 100
=====
task_id 75
Model answer: 8.0
Ground truth answer: 60.0
Correct: 5 out of 100
=====
task_id 76
Model answer: None
Ground truth answer: 5.0
Correct: 5 out of 100
=====
task_id 77
Model answer: None
Ground truth answer: 100.0
Correct: 5 out of 100
=====
task_id 78
Model answer: 27.0
Ground truth answer: 6.0
Correct: 5 out of 100
=====
task_id 79
Model answer: 70.0
Ground truth answer: 70.0
Correct: 6 out of 100
=====
task_id 80
```

```

Model answer: 14.0
Ground truth answer: 10.0
Correct: 6 out of 100
=====
task_id 81
Model answer: 33.0
Ground truth answer: 17.0
Correct: 6 out of 100
=====
task_id 82
Model answer: 624.0
Ground truth answer: 623.0
Correct: 6 out of 100
=====
task_id 83
Model answer: 1080.0
Ground truth answer: 600.0
Correct: 6 out of 100
=====
task_id 84
Model answer: 30.0
Ground truth answer: 15.0
Correct: 6 out of 100
=====
task_id 85
Model answer: None
Ground truth answer: 44.0
Correct: 6 out of 100
=====
task_id 86
Model answer: None
Ground truth answer: 22.0
Correct: 6 out of 100
=====
task_id 87
Model answer: None
Ground truth answer: 9360.0
Correct: 6 out of 100
=====
task_id 88
Model answer: 0.0
Ground truth answer: 8000.0
Correct: 6 out of 100
=====
task_id 89
Model answer: None
Ground truth answer: 24.0
Correct: 6 out of 100

```

```

=====
task_id 90
Model answer: None
Ground truth answer: 225.0
Correct: 6 out of 100
=====
task_id 91
Model answer: None
Ground truth answer: 28.0
Correct: 6 out of 100
=====
task_id 92
Model answer: None
Ground truth answer: 4.0
Correct: 6 out of 100
=====
task_id 93
Model answer: None
Ground truth answer: 36.0
Correct: 6 out of 100
=====
task_id 94
Model answer: None
Ground truth answer: 348.0
Correct: 6 out of 100
=====
task_id 95
Model answer: None
Ground truth answer: 40.0
Correct: 6 out of 100
=====
task_id 96
Model answer: 3.0
Ground truth answer: 3.0
Correct: 7 out of 100
=====
task_id 97
Model answer: 6.0
Ground truth answer: 12.0
Correct: 7 out of 100
=====
task_id 98
Model answer: 20.0
Ground truth answer: 5.0
Correct: 7 out of 100
=====
task_id 99
Model answer: 40.0

```


Ground truth answer: 58.0

Correct: 7 out of 100

=====

Final Accuracy: 0.01

```
[ ]: accuracy = all_correct / total* 100
      print(f"Final Accuracy Llama2 pretrained on trace: {accuracy:.2f}%")
```

Final Accuracy Llama2 pretrained on trace: 7.00%

```
[ ]: model_save_path = "model_save_path/llama-2-trace"
      model.save_pretrained(model_save_path)
```

```
[ ]:
```